

Vinaik Chhetri, Marc Salvadó Benasco

Abstract

This project aims to implement at least 3 architectures and recreate the results in Spektral. This paper discusses models in the following papers: “How Attentive are Graph Attention Networks?”, “Beyond Low-frequency Information in Graph Convolutional Networks” and “Simple and Deep Graph Convolutional Networks”. All three models were successfully implemented/ported and they were able to achieve similar results to the ones in the paper.

1 Introduction

A great amount of data can be modelled as a graph and GNNs (Graph Neural Networks) have been widely used to deal with such data. Spektral and PyG are Python libraries to train GNNs; the first one, based on Keras (TensorFlow), and the second one, on PyTorch.

Majority of the GNNs models are mainly implemented using PyG whereas Spektral lacks sufficient architectures/layers. Our task is then to implement some of these models in Spektral and recreate the results published in the papers. Two of the models that were implemented were only trained on synthetic datasets to see the core idea/principles behind the model. Successfully obtaining similar/same results on the synthetic datasets should guarantee that it would perform similarly on the other datasets the author experimented on.

The three models that were chosen were:

1. GATv2 (“How Attentive are Graph Attention Networks?”) [1]
2. FAGCN (“Beyond Low-frequency Information in Graph Convolutional Networks”) [2]
3. GCNII (“Simple and Deep Graph Convolutional Networks”) [3].

- GATv2

GATv2 is a natural extension of GAT(Graph Attention Model) which was engineered to overcome the static nature of GAT’s attention. Shaked Brody et.al. modify the existing GAT architecture to generate dynamic attention. To demonstrate the advantage of GATv2 over GAT, they propose/train on a synthetic dataset(“DICTIONARYLOOKUP”) and other datasets [4]. We simply tried to only recreate the results for the synthetic dataset as this dataset clearly showed GATv2’s superiority over GAT. The synthetic dataset generator has a parameter k which constructs a $k \times k$ bipartite graph dataset of cardinality $k!$. Shaked Brody et.al. trained their models over $k = 4...20$. We were successfully able to reconstruct the results but we could only train the models over $k = 4, 5$ due to high computation time.

- FAGCN

According to Deyu Boet et.al FAGCN is a generalization of most existing GNNs([2]). Most GNNs only make use of low frequency signals of node features. The low pass filters in GNNs are only good for assortative networks that always have similar nodes features connected making it unsuitable for inference on disassortative networks([2]). FAGCN was thus created to incorporate both types of signals during aggregation to deal with real world graphs that have a mix of assortative and disassortative edges/graphs. To demonstrate the advantage of FAGCN over low frequency filters and high frequency filters, Deyu Boet et.al proposes/trains on a synthetic dataset which has 2 classes of nodes with a parameter p and q to control the assortativity/disassortativity of the network [5]. We were able to successfully produce their results using their dataset.

- GCNII

GCNII is an extension of the well-known GCNs (Graph Convolutional Networks), which are a generalisation of CNN (Convolutional Neural Networks) for graph-structured data. GCNs are known to suffer from **oversmoothing**, which is the fact that, as the number of layers with non-linearities grows (increasing the depth of the model), the nodes outputs tend to be indistinguishable. This is why deep GCN perform worse than shallow GCN, which shallowness has the limitation of not being able to approximate the high-frequency components of the data. The goal of GCNII is then to provide more depth to GCN while avoiding the oversmoothing problem.

In section 2 we will introduce the main idea of the aforementioned papers; in section 3, we will explain the used methodology; in section 4, we will explain the steps that have been done during the implementation, as well as the main challenges and problems; in section 5, we will present a comparison of the results between our implementations and the existing implementations in PyG; and in section 6 we will discuss the results and conclude the paper, (as well as suggest future works).

2 Models

- GATv2

As mentioned above GATv2 is a modification of GAT with an enhanced attention mechanism. GAT, unlike other vanilla GNNs does not use an aggregation function that equally weighs all neighbours of a node but rather learns a vector of weights via a single layered neural network. This can be seen as follows:

$$\begin{aligned} e(h_i, h_j) &= \sigma_{lr}(a^T [Wh_i || Wh_j]) \\ \alpha_{ij} &= \frac{\exp(e(h_i, h_j))}{\sum_{j' \in N_i} \exp(e(h_i, h_{j'}))} \\ h'_i &= \sigma \left(\sum_{j \in N_i} \alpha_{ij} Wh_j \right) \end{aligned} \quad (1)$$

Note: σ_{lr} is LeakyReLU()

According to the author the computation of the unnormalized scores $e(h_i, h_j)$ creates static attention—Given a family of scoring functions $F \subseteq (\mathcal{R}^d \times \mathcal{R}^d \mapsto \mathcal{R})$ for a set of keys $\mathcal{K} = k_1, k_2, \dots, k_n \in \mathcal{R}^d$ and for a set of queries $\mathcal{Q} = q_1, q_2, \dots, q_m \in \mathcal{R}^d$, all functions $f \in F$ computes static attention if for each f there exists $j_f \in [n]$ such that for each query $i \in [m]$ and each key $j \in [n]$, $f(q_i, k_{j_f}) \geq f(q_i, k_j)$. In short no matter which function is chosen from a family of function that computes static attention, it will always have a key that scores highest regardless of the query. The author claims that the unnormalized scores $e(h_i, h_j)$ has both the parameters inside the leaky relu function thus constraining its degrees of freedom. Due to the monotonicity of leaky relu and softmax, $e(h_i, h_j)$ and the attention weights α_{ij} will always be maximum for maximum h^* among all hidden node representation thus generating static attention.

This problem inspired the authors to create GATv2 which computes dynamic attention—Given a family of scoring functions $F \subseteq (\mathcal{R}^d \times \mathcal{R}^d \mapsto \mathcal{R})$ for a set of keys $\mathcal{K} = k_1, k_2, \dots, k_n \in \mathcal{R}^d$ and for a set of queries $\mathcal{Q} = q_1, q_2, \dots, q_m \in \mathcal{R}^d$, all functions $f \in F$ computes dynamic attention if for any mapping $\phi : [m] \mapsto [n]$ there exists $f \in F$ such that for each query $i \in [m]$ and each key $j \neq \phi(i)$, $f(q_i, k_{\phi(i)}) \geq f(q_i, k_j)$. In shorts the maximal key is query dependent meaning that there exists a different key association for a different query. GATv2 fixes the problem by taking the weight vector a , seen in equation(1), outside of the LeakyRelu activation function thus giving us:

$$e(h_i, h_j) = a^T \sigma_{lr}(W[h_i || h_j]) \quad (2)$$

- FAGCN

The architecture of FAGCN involves first applying a multi-layered perceptron layer on the node features which is then sent through the adaptive frequency aggregation layer $L - 1$ times. Then a final affine transformation is applied to the output. The model can be seen as the following:

$$\begin{aligned} h_i^{(0)} &= \phi(W_1 h_i) \\ h_i^{(l)} &= \epsilon h_i^{(0)} + \sum_{j \in N_i} \frac{\alpha_{i,j}^G}{\sqrt{d_i d_j}} h_j^{(l-1)} \\ h_{out} &= W_2 h_i^{(L)} \\ \alpha_{ij}^G &= \tanh(g^T [h_i || h_j]) \end{aligned} \quad (3)$$

Note: $\alpha_{ij}^G = \alpha_{ij}^L - \alpha_{ij}^H$ where α_{ij}^L is the proportion of node j 's low frequency signal to node i and α_{ij}^H is the proportion of node j 's high frequency signal to node i . Since α_{ij}^G is an output of a tanh function, $\alpha_{ij}^L + \alpha_{ij}^H = 1$. If $\alpha_{ij}^G > 0$ then the low frequency signal from node j to i is greater than the high frequency signal from node j to i . g is a shared convolutional kernel between the connecting nodes. $\epsilon \in [0, 1]$.

- GCNII

Given a graph $\mathbb{G} = (\mathbb{V}, \mathbb{E})$, $|\mathbb{V}| = n$, $|\mathbb{E}| = m$; let's consider the same graph with a self loop per node, $\tilde{\mathbb{G}} = (\mathbb{V}, \tilde{\mathbb{E}})$, with adjacency matrix $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ and degree matrix $\tilde{\mathbf{D}} = \mathbf{D} + \mathbf{I}$, where \mathbf{A} and \mathbf{D} are the adjacency and degree matrices of \mathbb{G} .

Given the node features $\mathbf{X} \in \mathbb{R}^{n \times d}$, which dimension can be downsized via dense layers to $\mathbf{H} \in \mathbb{R}^{n \times \tilde{d}}$, $\tilde{d} < d$, a GCN layer is defined as $\mathbf{H}^{(l+1)} = \sigma(\tilde{\mathbf{P}} \mathbf{H}^{(l)} \mathbf{W}^{(l)})$, where \mathbf{H} are the hidden layers, \mathbf{W} are the weight matrices, σ stands for ReLU and $\tilde{\mathbf{P}} = \tilde{\mathbf{D}}^{-1/2} \cdot \tilde{\mathbf{A}} \cdot \tilde{\mathbf{D}}^{-1/2}$.

The GCNII layer suggests two improvements: initial residual connection (first term of the product inside the sigma), in order to ensure that each node output contains a fraction of the first hidden layer, despite the high number of convolutional layers; and identity mapping (second term of the product) —following the philosophy of ResNets, which have proven to work and have properties like having a small norm for \mathbf{W} —. Thus, the GCNII layer is defined as

$$\begin{aligned} \mathbf{H}^{(l+1)} &= \sigma \left(\left((1 - \alpha_l) \tilde{\mathbf{P}} \mathbf{H}^{(l)} + \alpha_l \mathbf{H}^{(0)} \right) \right. \\ &\quad \left. \left((1 - \beta_l) \mathbf{I}_n + \beta_l \mathbf{W}^{(l)} \right) \right) \end{aligned} \quad (4)$$

3 Methodology

The authors for GATv2 and FAGCN trained their models on multiple datasets but they both had synthetic datasets which demonstrated the core principle and idea behind their model so we only trained GATv2 and FAGCN on these datasets and recreated only results pertaining to them. And secondly due to time/experience and computational power constraints, we were not able to train the models on all the datasets. But we were able to train our third model GCNII on multiple datasets.

- GATv2

For this part of the project both GATv2 and GAT were implemented and trained on the synthetic dataset "DICTIONARYLOOKUP". The dataset is thoroughly explained in Section 4.1.. Firstly the models were only trained on datasets of 4! (4×4) bipartite graphs and 5! (5×5) bipartite graphs mainly due to inefficient training measures and also due to low computational power. Shaked Brody et.al have implemented a class that returns training and test dataset of $k \times k$ bipartite graphs. They implemented this using PyG and we made modifications to it using Spektral. They trained a single-headed and eight-headed GAT and GATv2 on the

”DICTIONARYLOOKUP” dataset for $k = 1, 2, \dots, 20$. They were able to show that a single-headed GATv2 outperforms an 8-headed and 1-headed GAT on training and test datasets. The only code that we used from the authors was the data generator. We used the implementation of GAT in Spektral as a starter. Our model only deals with a Single Loader meaning our model only takes in 1 graph at a time. This is clearly not optimal for training datasets with many graphs. Our model was implemented well but our training procedure is not optimal—the ideal situation would have been to use batch gradient descent but we instead averaged our gradients across a batch to make an update to the weights. This was done due to time constraints and lack of experience with Disjoint/Batch loaders and Spektral in general.

- FAGCN

The authors of FAGCN also used a synthetic dataset which is thoroughly described in section 4.1. They also implemented a data generator using numpy and SciPy which made it convenient for us to use but we still had to do some pre-processing and modifications for Spektral compatibility. The authors created 10 datasets each containing 200 nodes from 2 different classes and the datasets were created from low disassortativity to high disassortativity. A dataset of 200 nodes were split into 2—training and test dataset. These datasets were then fed into FAGCN to experiment with high, low and adaptive frequency filters. So the goal is to predict the labels for the nodes i.e. to determine which class the node belongs to. According to Deyu Boet et.al FAGCN can control the proportion of low/high frequency signals to node i from j via α_{ij}^G as shown in equation(3). The authors did not explicitly state how they trained on the dataset using high frequency and low frequency filters to produce Figure 1.a. in the paper ”Beyond Low-frequency Information in Graph Convolutional Networks” [2]. So we found it plausible to use $\alpha_{ij}^G > 0$ as a low frequency retainer/filter and $\alpha_{ij}^G < 0$ as a high frequency retainer/filter. Whereas α_{ij}^G in FAGCN is not strictly positive or negative as FAGCN has the property of an adaptive frequency filter. Hence in order to mimic a low frequency filter the α_{ij} matrix in our code is put in absolute to make all its entries positive. And in order to mimic a high frequency filter the α_{ij} matrix in our code is put in absolute value and multiplied by -1 to make all its entries negative. Whereas for FAGCN we leave it as it is. This approach of taking absolute value and negating does not purely retain high frequency signals and only taking absolute value of the α_{ij} matrix does not purely retain low frequency signals either. So a better alternative would be to set α_{ij} matrix to a matrix of ones for a pure low frequency signal filter and set α_{ij} matrix to a matrix of -1 for a pure high frequency signal filter. $1 = \alpha_{ij}^L + \alpha_{ij}^H$ so $\alpha_{ij}^L = 1$ would imply $\alpha_{ij}^H = 0$ and $\alpha_{ij}^H = 1$ would imply $\alpha_{ij}^L = 0$. However this version of the report does also consider pure low frequency and pure high frequency signal filters/models. So this version of the report considers 5 models—FAGCN, high frequency high proportion filter, low frequency high proportion filter, pure high frequency filter and pure low frequency filter. The names are self explanatory and the above explanations should definitely aid in understanding the names. The results in the appendix show that we get more similar results to that

of the authors when we consider pure high frequency filter and pure low frequency filter. Note that for a given dataset and model, the highest test accuracy for a model was chosen when the training loss was its minimum.

- GCNII

For the implementation in TensorFlow and Spektral, we took the original code [6] (in PyG) from the authors of the paper. The files that were modified were `train.py`, `model.py` and `utils.py`. The main challenge was to learn how to adapt the functions and operations from PyTorch to TensorFlow and Spektral; some of them couldn’t be ported from one library to another, but they had to be rewritten in a different way, like the dataloaders and the training process.

A problem we encountered was that we couldn’t apply the Weight-Decay regularization to the Adam optimizer the same way it was implemented in the original code; that is, separately for the dense layers and for the convolutional layers. Instead, we have implemented it jointly for all kind of layers.

4 Implementation

4.1 Datasets

- GATv2

Authors of GATv2 propose a synthetic dataset (DICTIONARYLOOKUP). They build a data generator that accepts parameter k and builds a $k! \times k$ bipartite graph with 2 attributes each for all nodes. The first $1, \dots, k$ nodes are the query nodes and the second set of k nodes $k+1, \dots, 2k$ are key nodes. The 2nd attribute for the nodes are values whereas as the 1st attribute are the name/identifier of the node. For the first set of nodes/queries $1, \dots, k$, the second attributes are unknown. Based on its 1st attribute and edge connections it needs to determine its value attribute.

- FAGCN

The authors propose a synthetic dataset and they built a data generator for it. The authors create 10 synthetic datasets ranging from small disassortativity to high i.e. $q = 0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.09, 0.1$. Each dataset has 200 nodes and each node attributes are 20 dimensional vectors. 100 nodes belong to class 1 and the other 100 belong to another class. Each entries in the node attribute vectors of class 1 and class 2 are sampled from $\mathcal{N}(0.5, 1)$ and $\mathcal{N}(-0.5, 1)$ respectively. This is done in order to make node attributes from same class similar to each other and node attributes from different classes dissimilar from one another. The data generator takes in parameters p and q . p determines the probability for a node in a class to connect to a node in the same class whereas q determines the probability for a node in a class to connect to a node in another class. p is fixed to 0.05 and q ranges from 0.01 to 0.1 with an increment of 0.01. According to Deyu Bo et.al ”the main idea is to gradually increase the disassortativity of the synthetic networks, and observe how the performance of these two signals changes”. Given an arbitrary dataset D with $p = p'$ and $q = q'$, D is split into 2 equal halves such that the training set and test set each have 50 nodes from class 1

and 50 nodes from class 2. The network is now trained over all the nodes and edges, however the targets/labels of nodes from the test set are masked.

- GCNII

The original code provides 12 different 1-graph-datasets, some of them are very large and take a lot of time to train. We have selected 4 of them, called “pubmed”, “cornell”, “texas” and “wisconsin”. “Pubmed” is a benchmark dataset about citations, taken from [7], where nodes are documents; edges, citations; and node features, a bag-of-words representation of the document. The other three datasets are web networks, where nodes are web pages; edges, hyperlinks; and node features, again, a bag-of-words representation of the page; and which are taken from [8].

The number of nodes of the datasets are, respectively: 19717, 183, 183 and 251; the number of edges, respectively: 44338, 295, 309 and 499; the dimension of the node features, respectively: 500, 1703, 1703 and 1703; and their number of classes, respectively: 3, 5, 5 and 5.

4.2 Hyperparameters

- GATv2

Dropout was set to 0 to see GAT and GATv2’s ability to generalize. Learning rate was set to 0.001 as the authors chose it as default. Experiments were conducted using 8 heads and 1 heads.

- FAGCN

Learning rate was set to 0.01 and $\epsilon = 0.3$ because the authors chose it as default. In addition taking additional 2 hyper-parameters into account for hyper-parameter tuning would be computationally expensive. We also only had 1 frequency aggregation layer. The authors did not discuss about their hyper-parameters concerning the synthetic dataset so we decided to perform hyper-parameter tuning over the following hidden dimensions and dropout rates : $(h, d) = \{32, 64, 128\} \times \{0.1, 0.2, 0.3, 0.4, 0.5\}$ where h :hidden dimension and d :dropout rate. We performed hyper-paramter tuning over 5 models— FAGCN, pure low pass filter, pure high pass filter, high proportion low frequency filter and high proportion high frequency filter. Initially in our report we only used FAGCN, high proportion low frequency filter and high proportion high frequency filter so for the final report we are also including pure low pass filter and pure high pass filter models. And in actuality the authors result/plot only concerns pure low pass filter and pure high pass filter models so we decided to choose the plot of FAGCN, pure low pass filter and pure high pass filter as our final plot to compare with the author’s. The graphs concerning high proportion low frequency filter and high proportion high frequency filter models can be found in the appendix. **In order to see the results of the hyper-parameter tuning we request you too look at the Appendix page.** Out of the graphs that were constructed in the Appendix page we tried to find the best match to the authors graphs thus giving us Figure 2.

- GCNII

As mentioned above, we couldn’t use the Adam with weight decay as in the original code, so we used tensorflow-addons’ AdamW with a unique weight-decay coefficient for all kind of layers. To find the best learning rate or weight-decay, we tried the values of 0.1, 0.01 and 0.001; for weight-decay, $1e-2$, $1e-3$, $5e-4$ and $1e-4$. In general, we started with the same parameters as in the original code, and from that finding settings with a better performance.

The rest of the hyperparameters —number of epochs, patience before early stopping, number of hidden layers in the fully-connected layers, dropout, α and λ , are the same as in the original code, in order to recreate at best the original model.

One discrepancy between our code and the original one is the fact that we have applied the convolutional formula as in the paper (Equation 4), whereas the original code computes A instead of \tilde{P} .

4.3 Experimental setup

- GATv2

GATv2 was trained on a free google Colab account. As described in section 3, GATv2 was not trained in an optimal way—all its operations were serialized, hence using a GPU did not help and at times it did deter the performance. So this model was strictly trained on a CPU(Intel(R) Xeon(R) CPU @ 2.20GHz). It used 1.44GB/12.68GB RAM and 38.36GB/107.72GB disk space. Link to the colab notebook: <https://colab.research.google.com/drive/1uNi8zvNVZFgfbf8553YbwiV7lzJdwDUaM?usp=sharing>

- FAGCN

FAGCN was also trained on a free google Colab account. This task was not computationally expensive due to its small dataset size so it was entirely trained on a CPU(Intel(R) Xeon(R) CPU @ 2.20GHz). It used 1.44GB/12.68GB RAM and 38.36GB/107.72GB disk space. Link to the colab notebook: <https://colab.research.google.com/drive/1ZkCK6rolHv6NU6HOht-ZH9fZGinVzSY?usp=sharing>

- GCNII

The experiments were run using a Google Colaboratory notebook and with the GPU enabled. Link to notebook with instructions: <https://colab.research.google.com/drive/1gqyNXnt3x4qtS9ozcLqEzRaLLE4WxzVJ?usp=sharing>.

4.4 Computational cost

- GATv2

Table 1. .

Bipartite graph size(kxk)	4x4	5x5
GAT-1head	26'	8h+
GATv2-1head	35'	3h 59'
GAT-8heads	20'	1h 50'
GATv2-8heads	32'	2h

- FAGCN

Training FAGCN on a synthetic dataset was very fast. On an average each training phase only took 18 seconds. But when performing hyper-parameter tuning over 15 hyper-parameter tuples (hidden dimension, dropout rate) for each of the 5 models, it took approximately 3 hours. During the tuning, the test accuracy for the lowest training/validation loss was saved. The saved checkpoint will be provided along with the code in order to reproduce results in the Appendix.

- GCNII

Computational power: an NVIDIA Tesla K80" GPU (the one from Google Colaboratory). The running times for each dataset are shown in Table 1.

Table 2. Running time of GCNII-Spektral per dataset.

Datasets	Running time (s)
Pubm	2698.1
Corn	12.4
Texa	35.7
Wisc	12.5

5 Results

- GATv2

Table 3. Training accuracy

Bipartite graph size(kxk)	4x4	5x5
GAT-1head	100	-
GATv2-1head	100	100
GAT-8heads	100	100
GATv2-8heads	100	100

Table 4. Testing accuracy

Bipartite graph size(kxk)	4x4	5x5
GAT-1head	90	-
GATv2-1head	100	100
GAT-8heads	95	90
GATv2-8heads	100	96.6

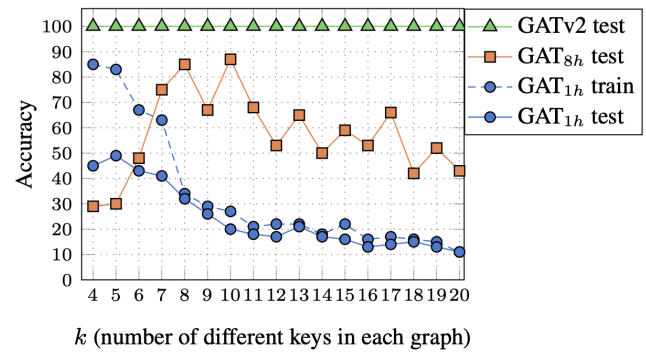


Figure 1. Author's plot ([1])

- FAGCN




Figure 2. One of our accuracy plots using hyper-parameters $(h, d) = (128, 0.3)$




Figure 3. Author's accuracy plot([2])

Table 5. Our Test accuracy FAGCN

Filters \ q	0.01	0.02	0.03	0.04	0.05
Pure LowFreq	0.95	0.88	0.76	0.67	0.6
Pure HighFreq	0.58	0.44	0.4	0.41	0.56
FAGCN	0.98	0.86	0.83	0.86	0.89

Filters \ q	0.06	0.07	0.08	0.09	0.10
Pure LowFreq	0.41	0.42	0.42	0.33	0.39
Pure HighFreq	0.58	0.75	0.71	0.9	0.94
FAGCN	0.87	0.86	0.79	0.86	0.95

- GCNII

In Table 6, one can observe the comparison between the test accuracy of the implementation from the authors of the paper (using PyG) and the one from our implementation (using Spektral).

Table 6. Comparison of the test accuracy between the expected accuracy (the one from the implementation in PyG) and the one from our implementation (Spektral).

Datasets	Acc (PyG)	Acc (Spektral)
Pubm	80.3	58.4
Corn	76.5	73.0
Texa	77.8	67.6
Wisc	81.6	68.6

The only dataset which results are comparable to the original code is Cornell (difference of 3.5%). The rest present a significant gap (between 10 and 20%) between the expected and the obtained accuracies. We attribute this to the fact that we couldn't apply the weight decay (regularisation) properly, and to the aforementioned discrepancy between A and \tilde{P} . Finally, the difference between validation and test accuracy for the dataset *Wisconsin* suggests us that either there is a problem of small data or that the train-validation-test partition indices are somehow biased, but anyway the choice of the partition makes a considerable difference in performance.

The evolution of the validation accuracy during the training process for the 4 datasets can be found in Figure 4, Figure 5, Figure 6 and Figure 7.




Figure 4. Validation accuracy evolution of our implementation of GCNII for the Pubmed dataset.




Figure 5. Validation accuracy evolution of our implementation of GCNII for the Cornell dataset.




Figure 6. Validation accuracy evolution of our implementation of GCNII for the Texas dataset.




Figure 7. Validation accuracy evolution of our implementation of GCNII for the Wisconsin dataset.

6 Discussion and conclusion

We conclude that the some of the implementations have gotten results as good as expected or very similar. For example, GATv2 and FAGCN have successfully replicated the models from the papers to Tensorflow-Spektral.

Table 3 and Table 4 show results for only 4×4 and 5×5 bipartite graphs so our comparison can only be made for the $k = 4$ and $k = 5$ datasets. Clearly our GATv2 with 1 head produces the exact results. GATv2 with 8 heads was not experimented by the authors but we nevertheless experimented with it to check for consistency that adding more heads improves performance. Our GAT with 1 head for a 4×4 bipartite graph dataset is very similar to the author's one if we consider its training accuracy, however our test accuracy is far higher than theirs. Our GAT with 1 head for a 5×5 bipartite graph dataset is left empty in Table 3 and Table 4 because the training was too slow. The test accuracy for GAT with 8 heads is clearly far different for us and the author. Despite certain setbacks we were still able to reproduce the exact results for GATv2. But it definitely would have been much better if we could have trained the models for higher k 's as we can see from Figure 1 that the real differences in performances occur as k increases. Future work concerning GATv2 would be to improve the training procedure by utilising batch gradient descent via Disjoint Loader/Batch Loaders. Experimenting with other datasets of the authors would also be a future work.

Figure 2 and Figure 3 show that the our and the author's plots are very similar. In our case the pure high pass filter model gives slightly high test accuracy at the start ($q = 0.01$) which makes our curve non-monotonic. This maybe the only

difference. However Figure 2 still shows that FAGCN is consistent/robust across values of q 's, pure high pass filter's test accuracy increases as disassortativity increases and pure low pass filter's test accuracy decreases as disassortativity increases. **More plots can be found in the appendix and maybe even a better match to the author's plot.** All the plots in appendix are actually quite similar to Figure 3 in the sense that the 3 curves have similar trends and order. If we observe Figure 3, we see that the author's curves have colorful shades. These shades are presumably multiple curves across different datasets with the same q . The dark lines in the shades are possibly the average in the shadings. Observe that FAGCN has the lowest variance and the other 2 curves have higher variance. We did not produce the shadings due to time constraints and this can be possibly a future work. Deyu Bo et.al experimented with different well-known assortative and disassortative datasets, however we only experimented with the synthetic dataset. The future work can possibly also entail training on these other datasets. We believe that the FAGCN model is ready to be committed/published to the Spektral library.

However, the implementation of GCNII has run into challenges that have significantly reduced their performance. As written above, we attribute the discrepancies between the expected results and the ones obtained to the inability to implement the Weight Decay regularisation as in the original paper. *Update with respect to the previous submission: we have rewritten the model script, a bit less similar to the original code, and more to the original paper —this didn't have a significant effect on the performance—. Besides, we have implemented AdamW regularization; even though it is not done as in the original paper (separately for convolutional and linear layers), it has shown a significant improvement with respect to Adam without regularization, with an improvement of 10% in Texas and Wisconsin, 3% on Cornell and a decay of 2% on Pubmed.

References

- [1] Shaked Brody, Uri Alon, and Eran Yahav. How attentive are graph attention networks? In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=F72ximsx7C1>.
- [2] Deyu Bo, Xiao Wang, Chuan Shi, and Huawei Shen. Beyond low-frequency information in graph convolutional networks. In *AAAI*. AAAI Press, 2021.
- [3] Zhewei Wei Ming Chen, Bolin Ding Zengfeng Huang, and Yaliang Li. Simple and deep graph convolutional networks. 2020.
- [4] (User) tech srl. Github — how attentive are gats. https://github.com/tech-srl/how_attentive_are_gats, 2022.
- [5] (User) bdy9527. Github — FAGCN. <https://github.com/bdy9527/FAGCN>, 2021.
- [6] Ming Chen (chennnM). Github — GCNII: Simple and Deep Graph Convolutional Networks. <https://github.com/chennnM/GCNII/tree/ca91f5686c4cd09cc1c6f98431a5d5b7e36acc92>, 2020. [Online; accessed May-2022].
- [7] Namata G. Bilgic M. Getoor L. Galligher B. Eliassi-Rad T. Sen, P. Collective classification in network data. *ai magazine*, 29(3), 93. 2008.
- [8] (User) graphdml-uiuc jlu. Github — Geom-GCN: Geometric Graph Convolutional Networks. <https://github.com/graphdml-uiuc-jlu/geom-gcn>, 2021.

7 Appendix

Accuracy plots of FAGCN, pure low pass filter and pure high pass filter

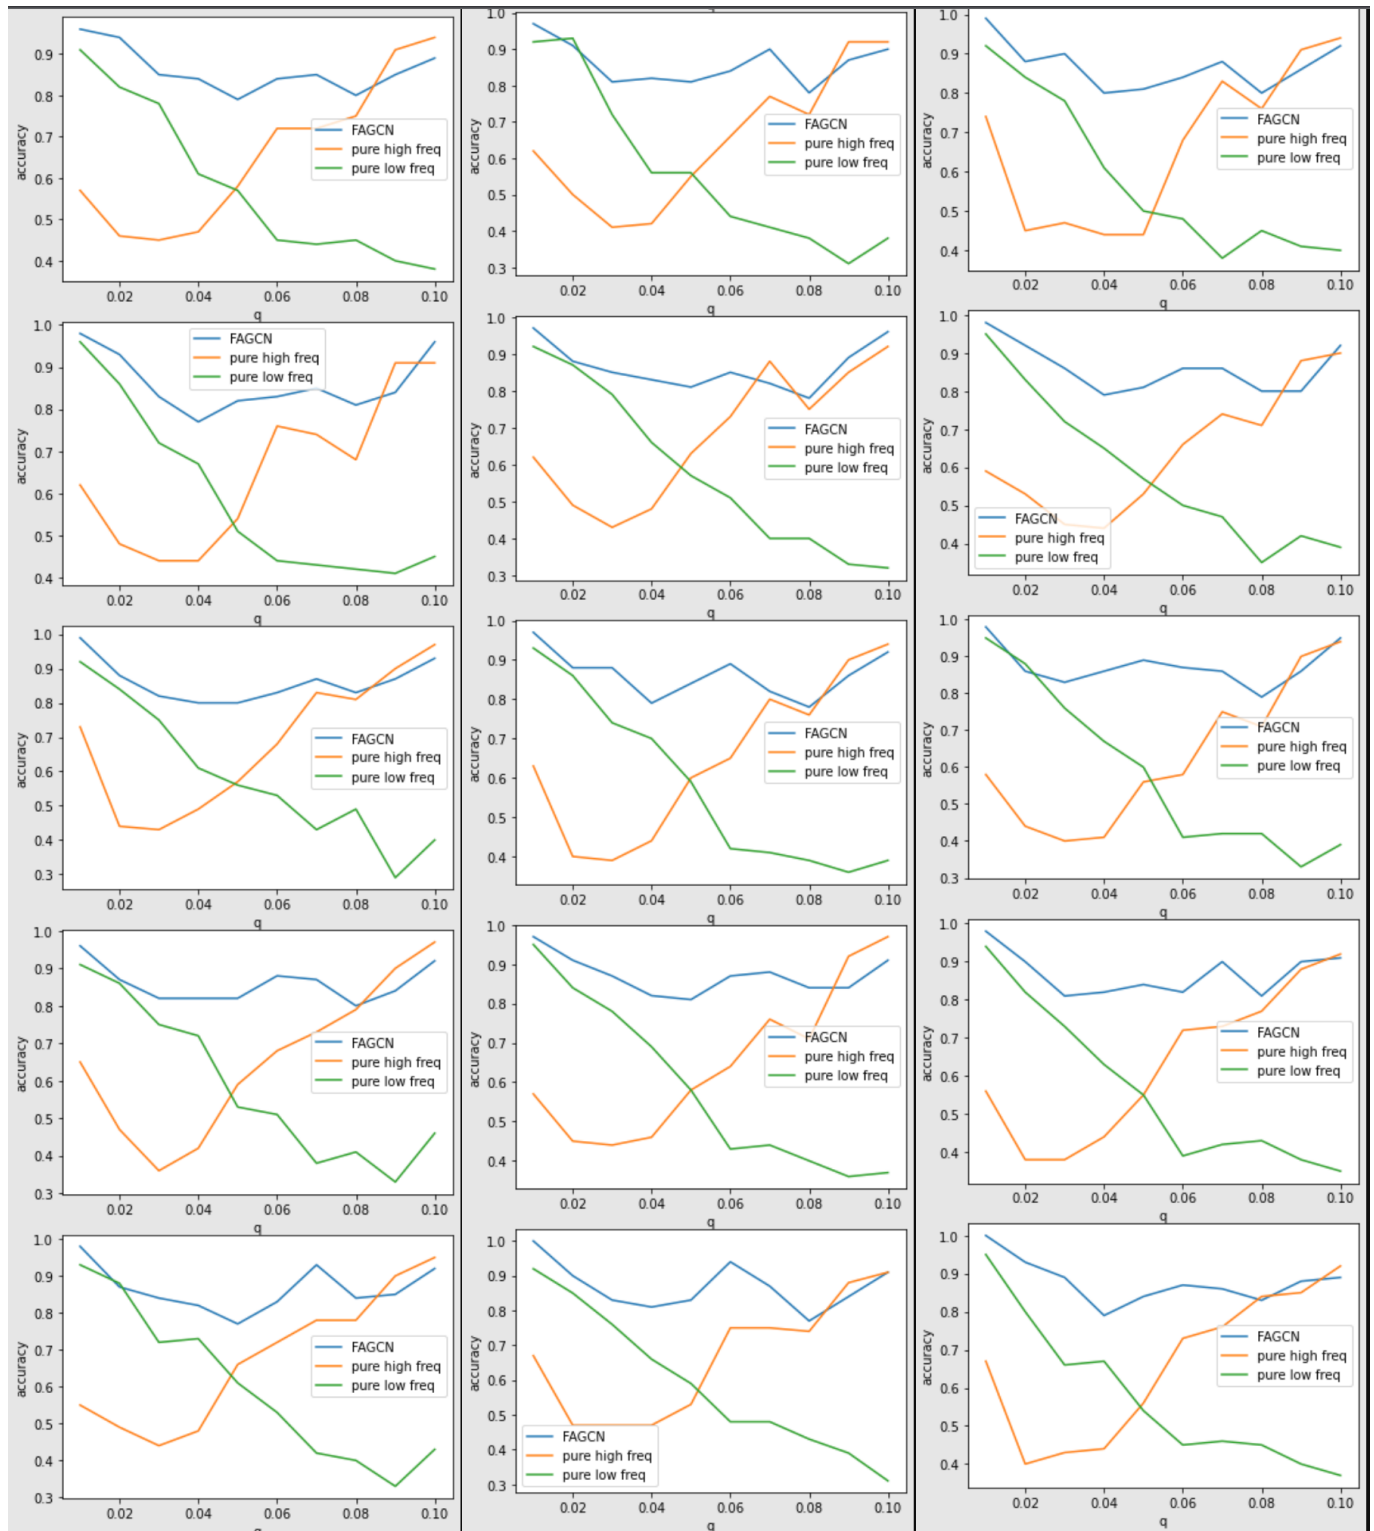


Figure 8. Accuracy plots using hyper-parameters $(h, d) = \{32, 64, 128\} \times \{0.1, 0.2, 0.3, 0.4, 0.5\}$

Note: The plots in the first, second and third columns correspond to hidden dimension $\{32, 64, 128\}$ and the rows correspond to the dropout rate $\{0.1, 0.2, 0.3, 0.4, 0.5\}$.

Accuracy plots of FAGCN, high proportion low frequency filter and high proportion high frequency filter

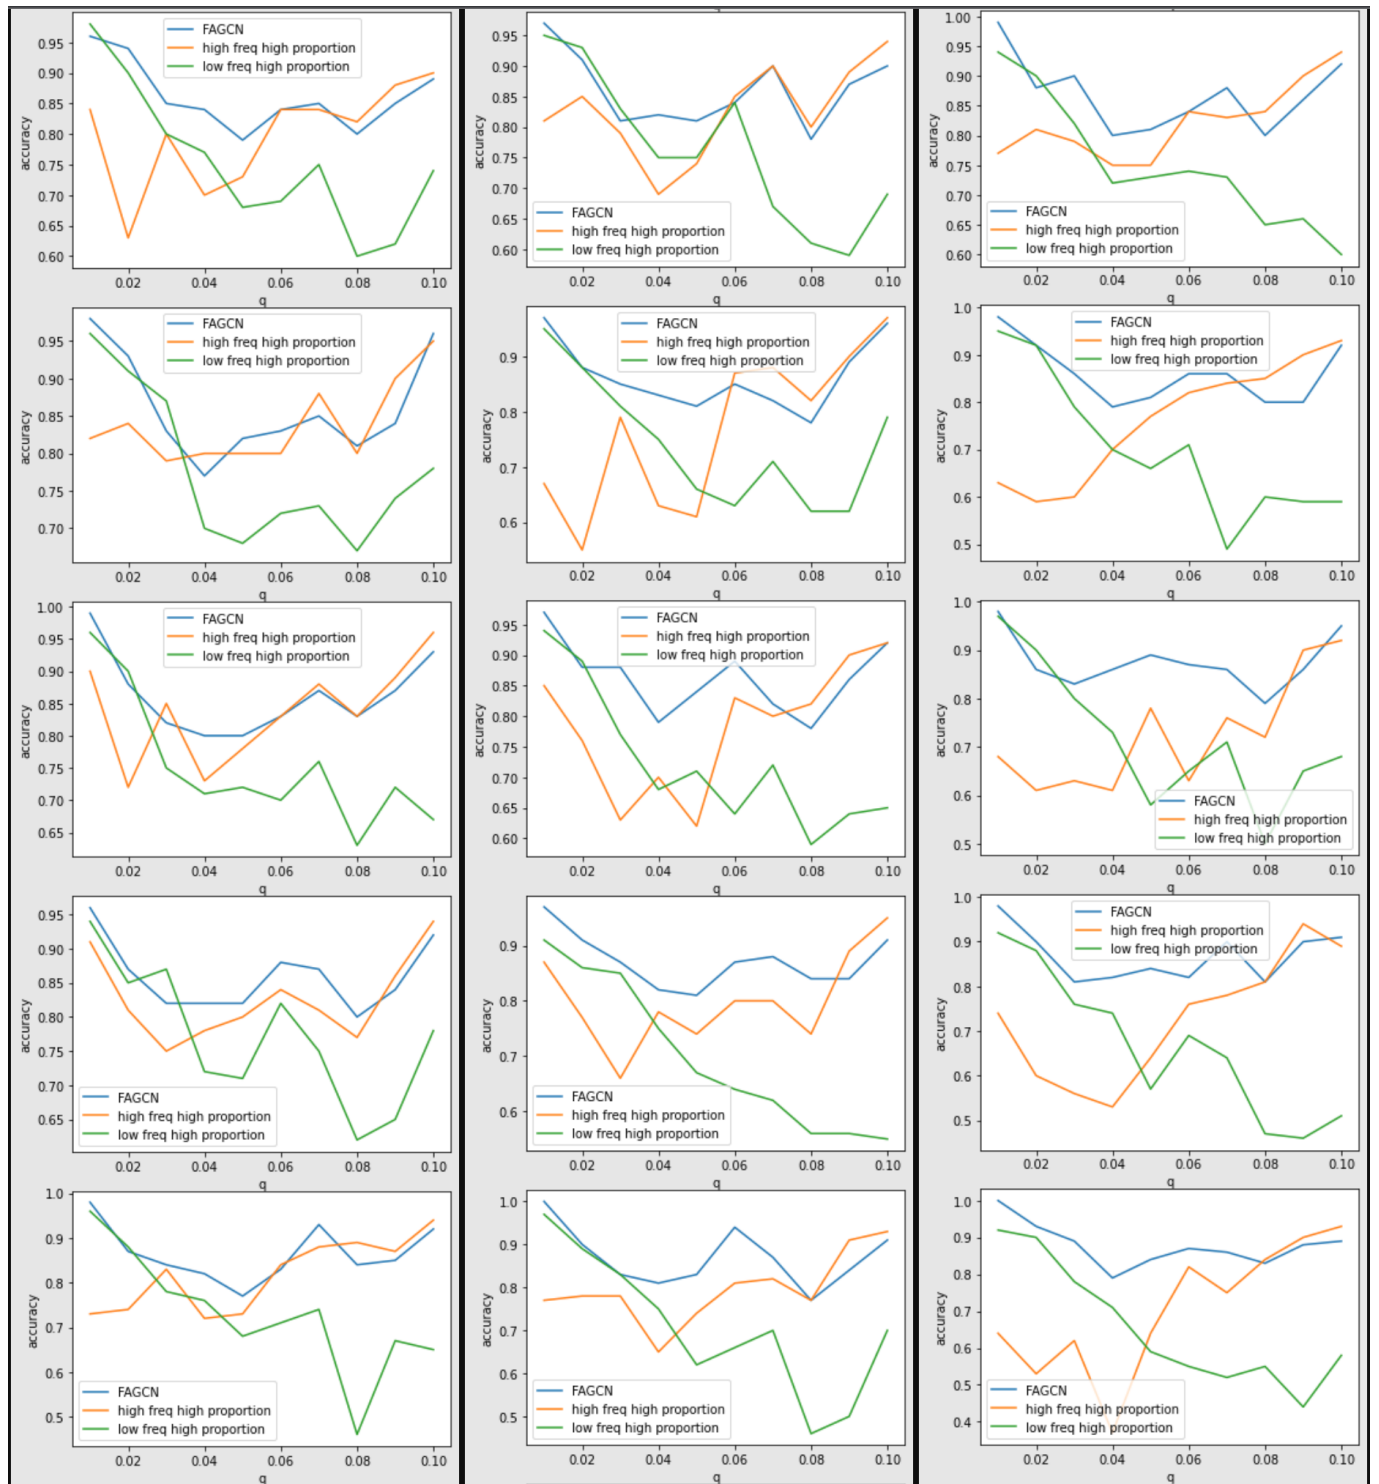


Figure 9. Accuracy plots using hyper-parameters $(h, d) = \{32, 64, 128\} \times \{0.1, 0.2, 0.3, 0.4, 0.5\}$

Note: The plots in the first, second and third columns correspond to hidden dimension $\{32, 64, 128\}$ and the rows correspond to the dropout rate $\{0.1, 0.2, 0.3, 0.4, 0.5\}$.